

@codecentric

Kafka 101_

A Messaging Service

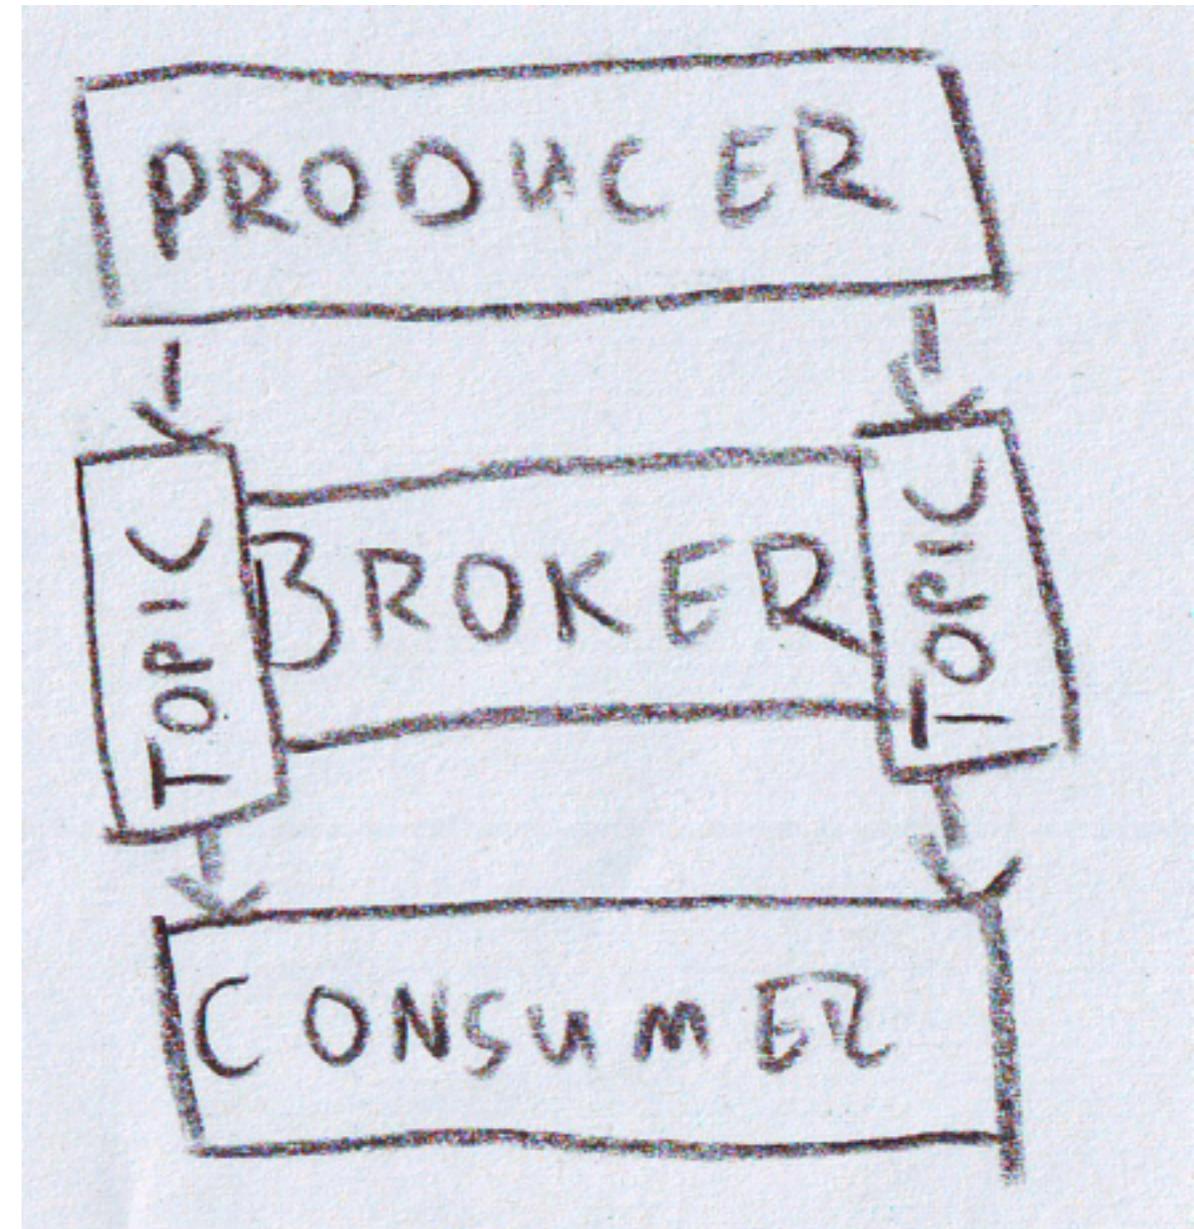
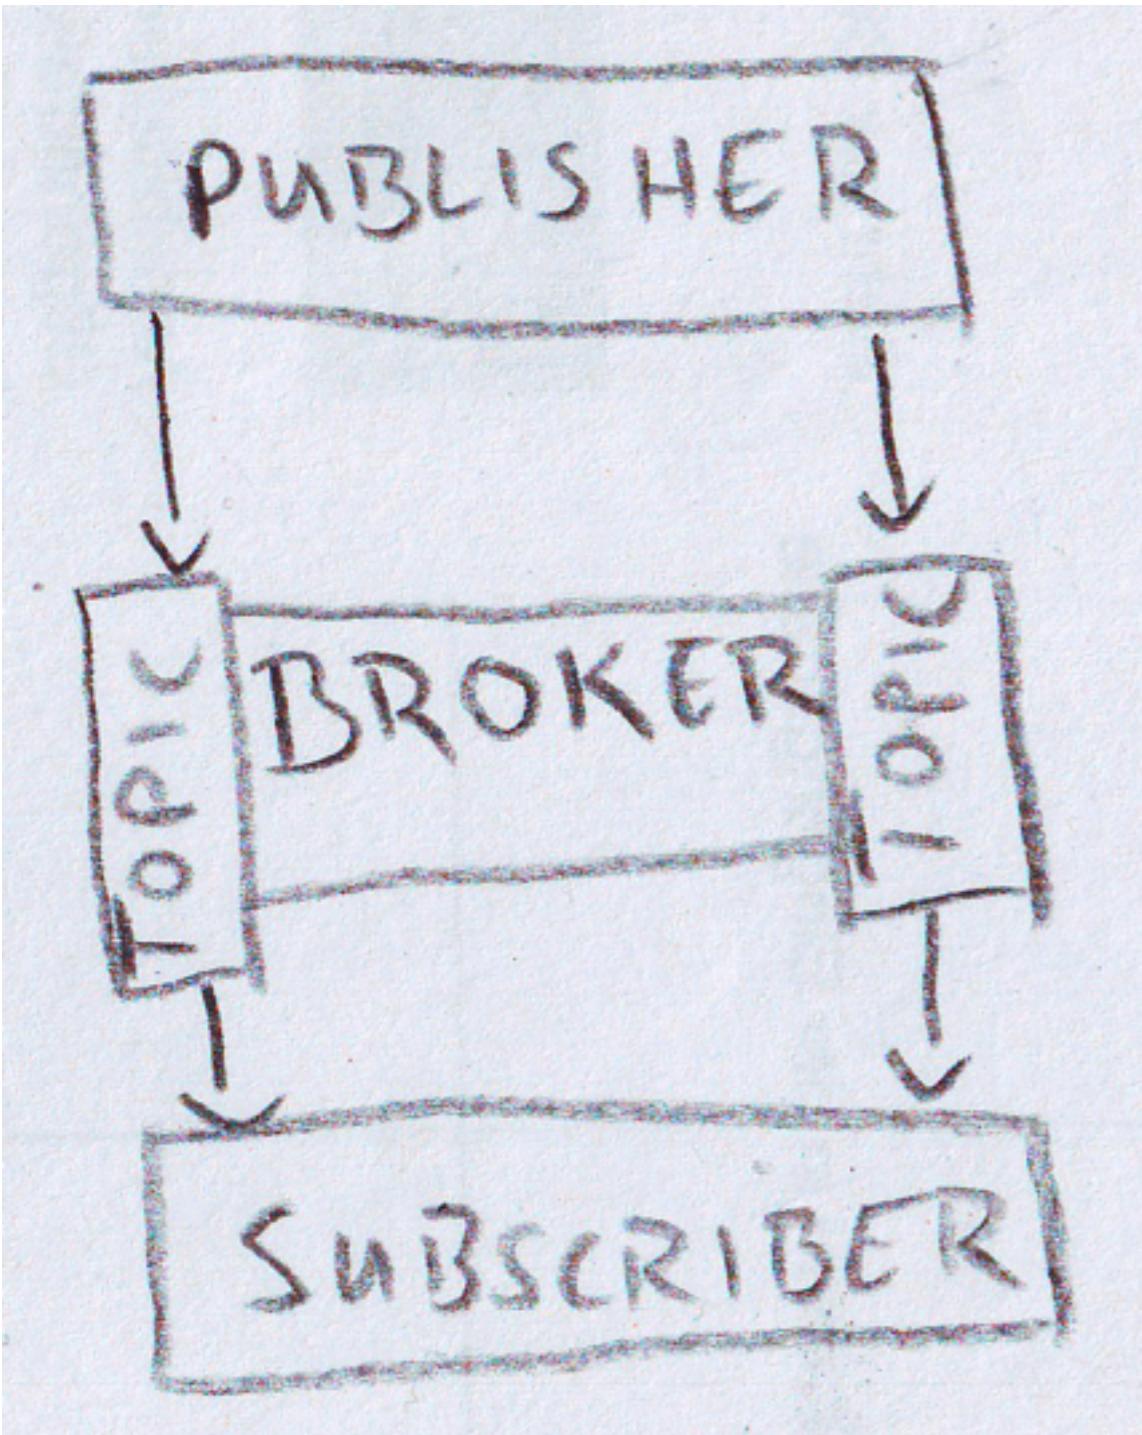
- Pub-Sub
- Distributed
- Replicated
- Commit Log based
- High Throughput
- Scalable



Where did it come from?

- Developed at LinkedIn
- Open sourced in early 2011 at ASF
- Graduated from incubator in October 2012
- Original Developers formed Confluent in 2014

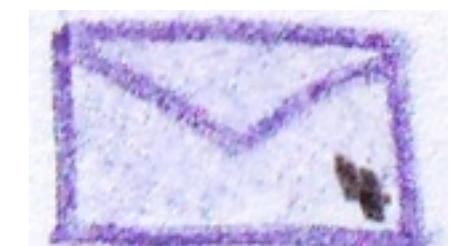
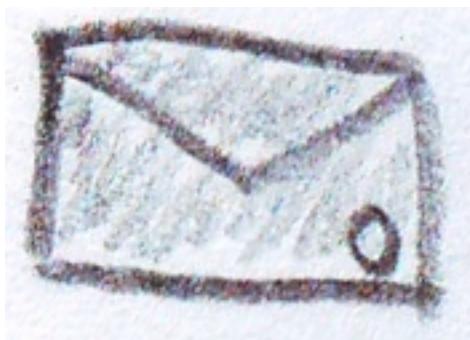
Pub-Sub Messaging



Messages

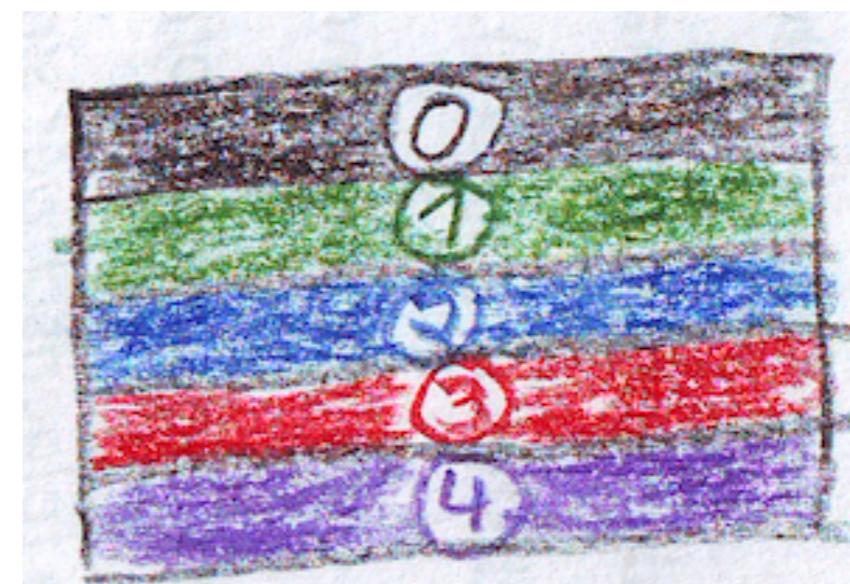
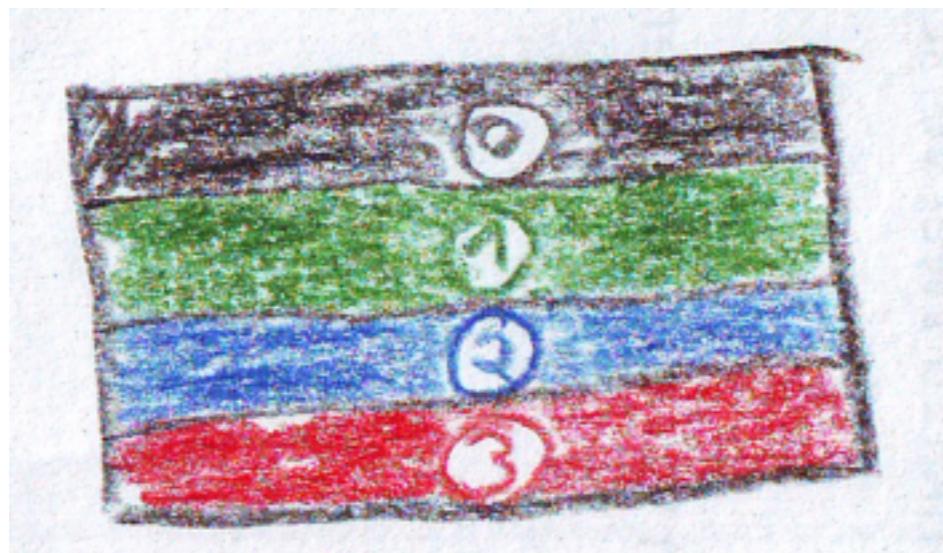
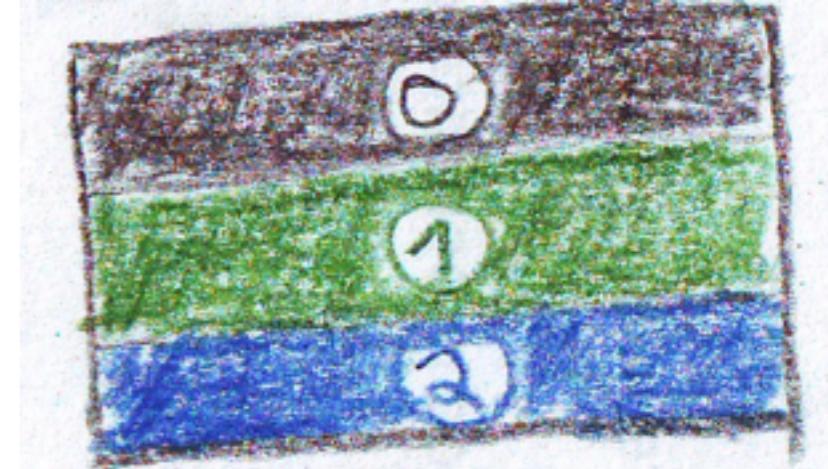
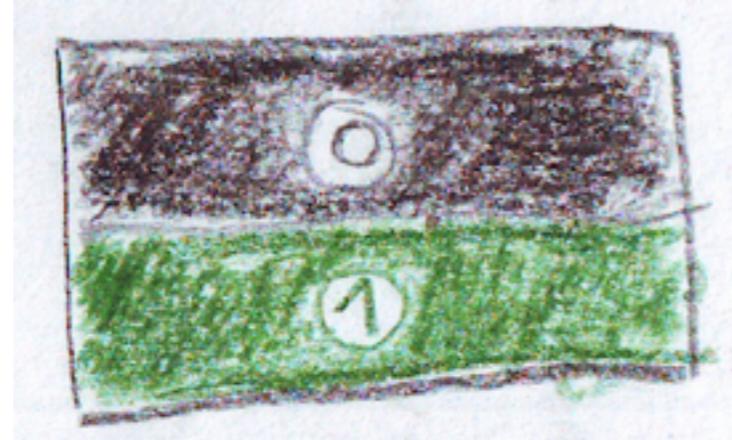
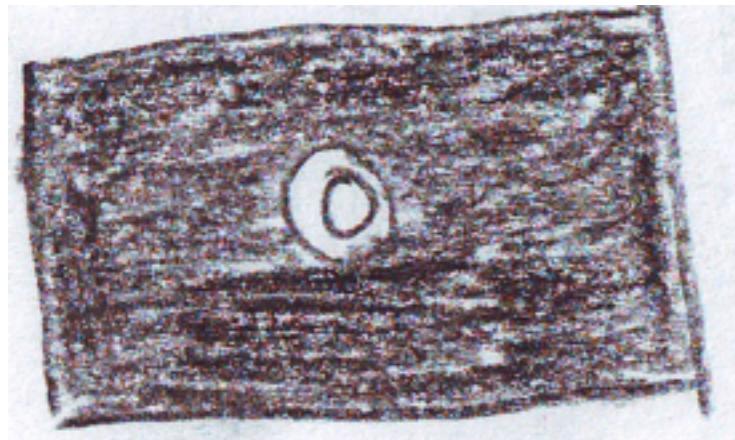
Key-Value Pairs

- Just bytes
- Can be zipped
- Arbitrary size
- Keys not necessarily unique

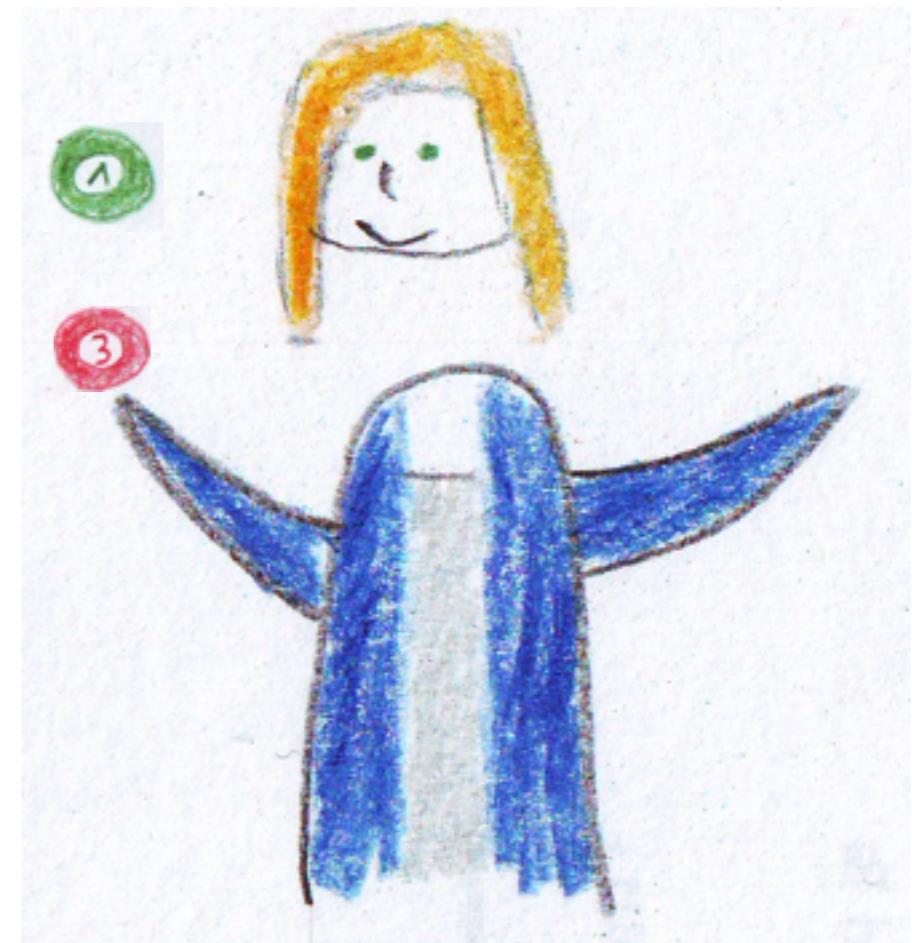


Topics, Partitions and Brokers

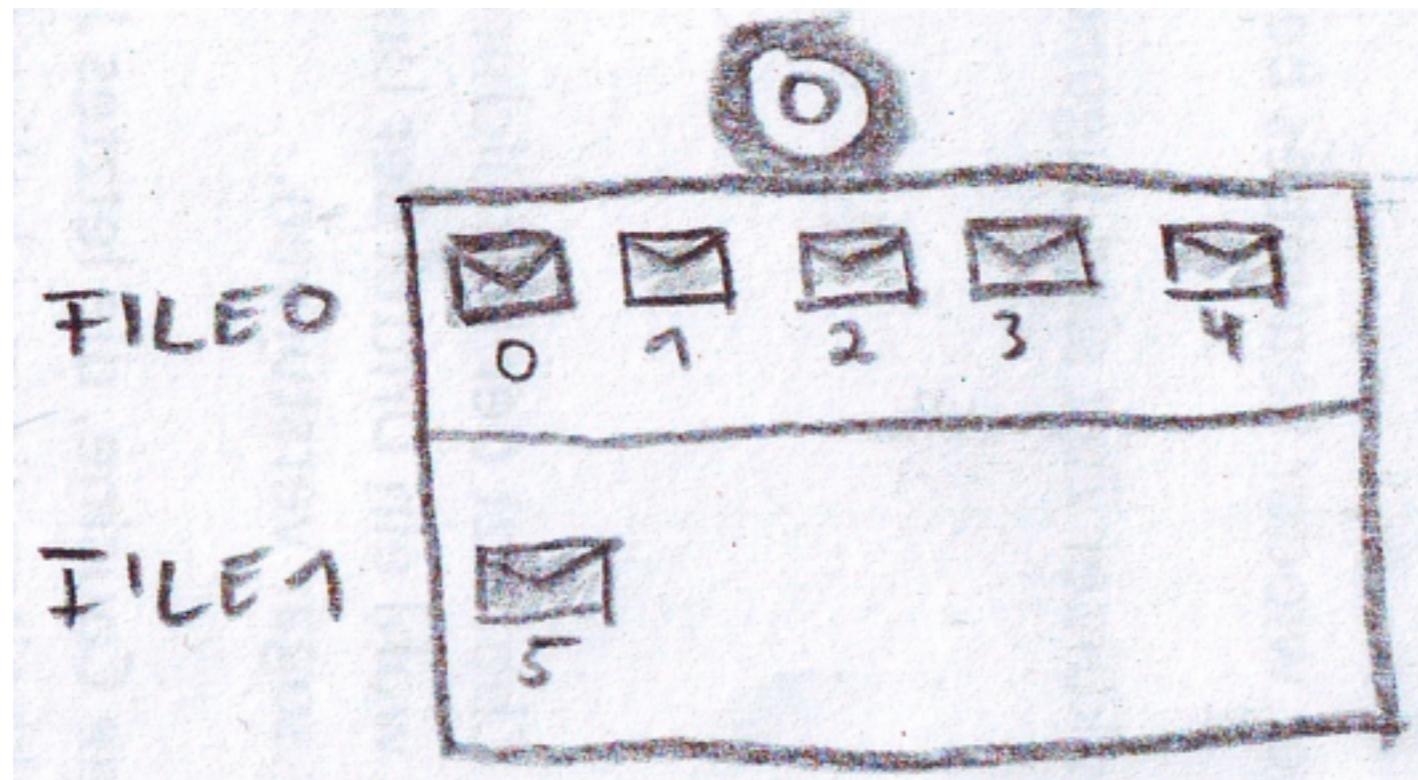
Topics are partitioned!



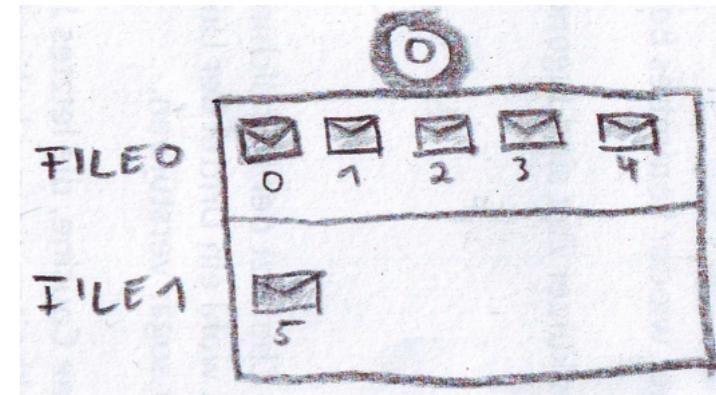
Partitions are assigned to Brokers



Partitions are stored as a commit log

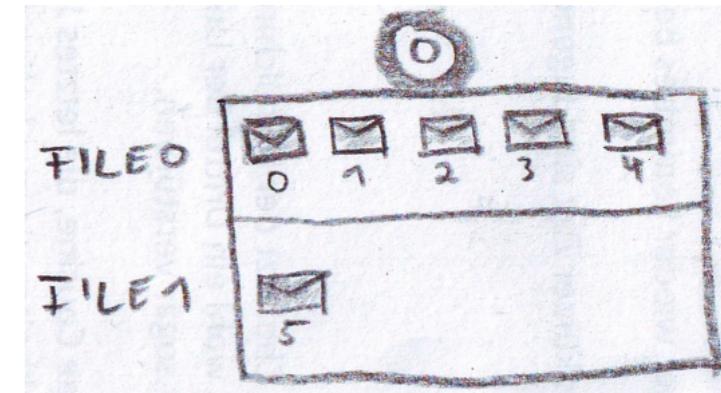


Messages are appended to partitions



- **n partitions per Broker (even from single topic!)**
- **Each message in a partition gets a unique offset**
 - **uniqueness per partition!**

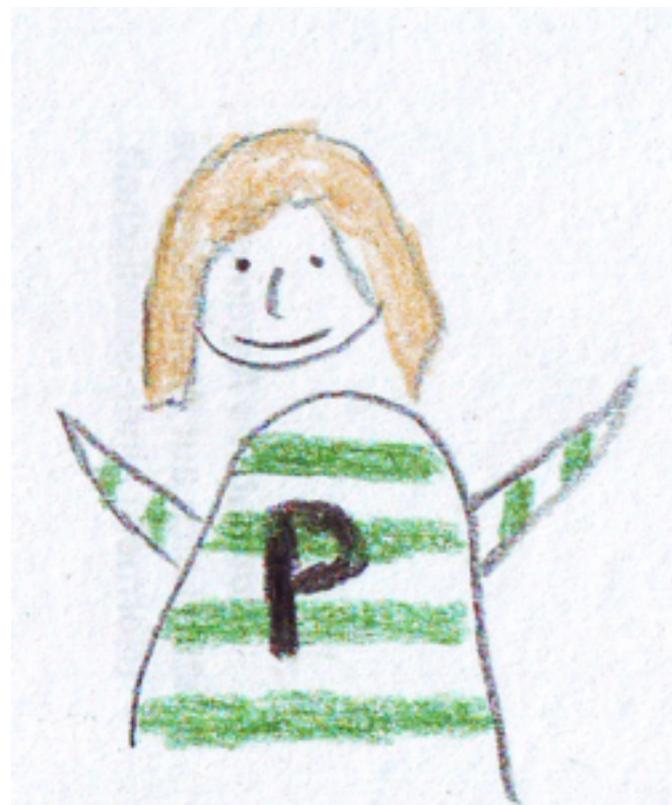
Don't fear the file system!



- Append-only files (1..n per partitions)
- Configurable retention time
- Optimised for linear reads and writes
- In most use cases no disk access due to page cache
- Zero-Copy
 - Data goes directly from page cache to socket!
 - allows Kafka to have (comparatively) low memory footprint
 - Not for SSL :(

Producing messages

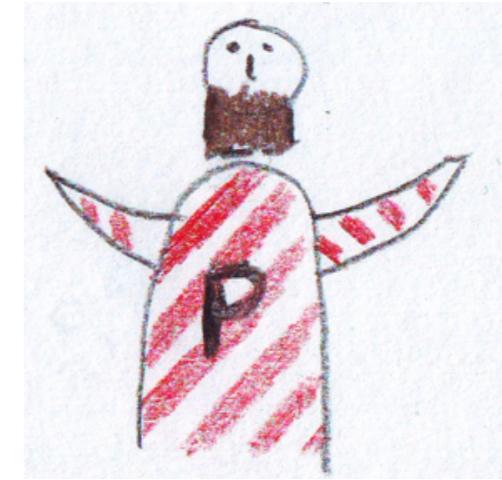
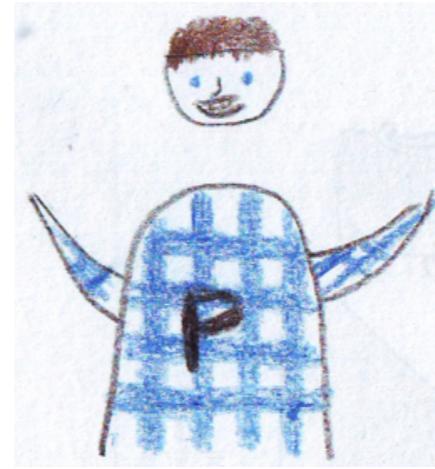
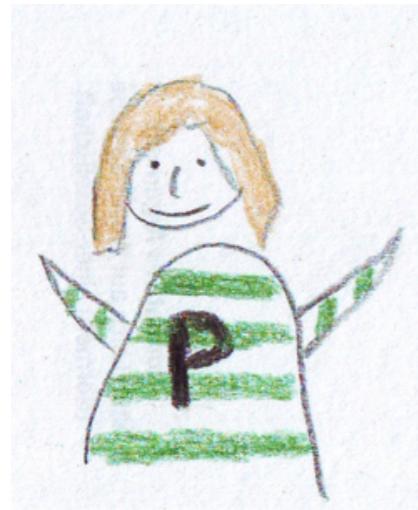
Meet the Producers



Producers send Messages to Brokers

- They decide which partition a message belongs to
- Messages always get sent directly to responsible Broker
- Producers specify expected consistency levels

Producing Messages



Consuming messages

Hello Consumers



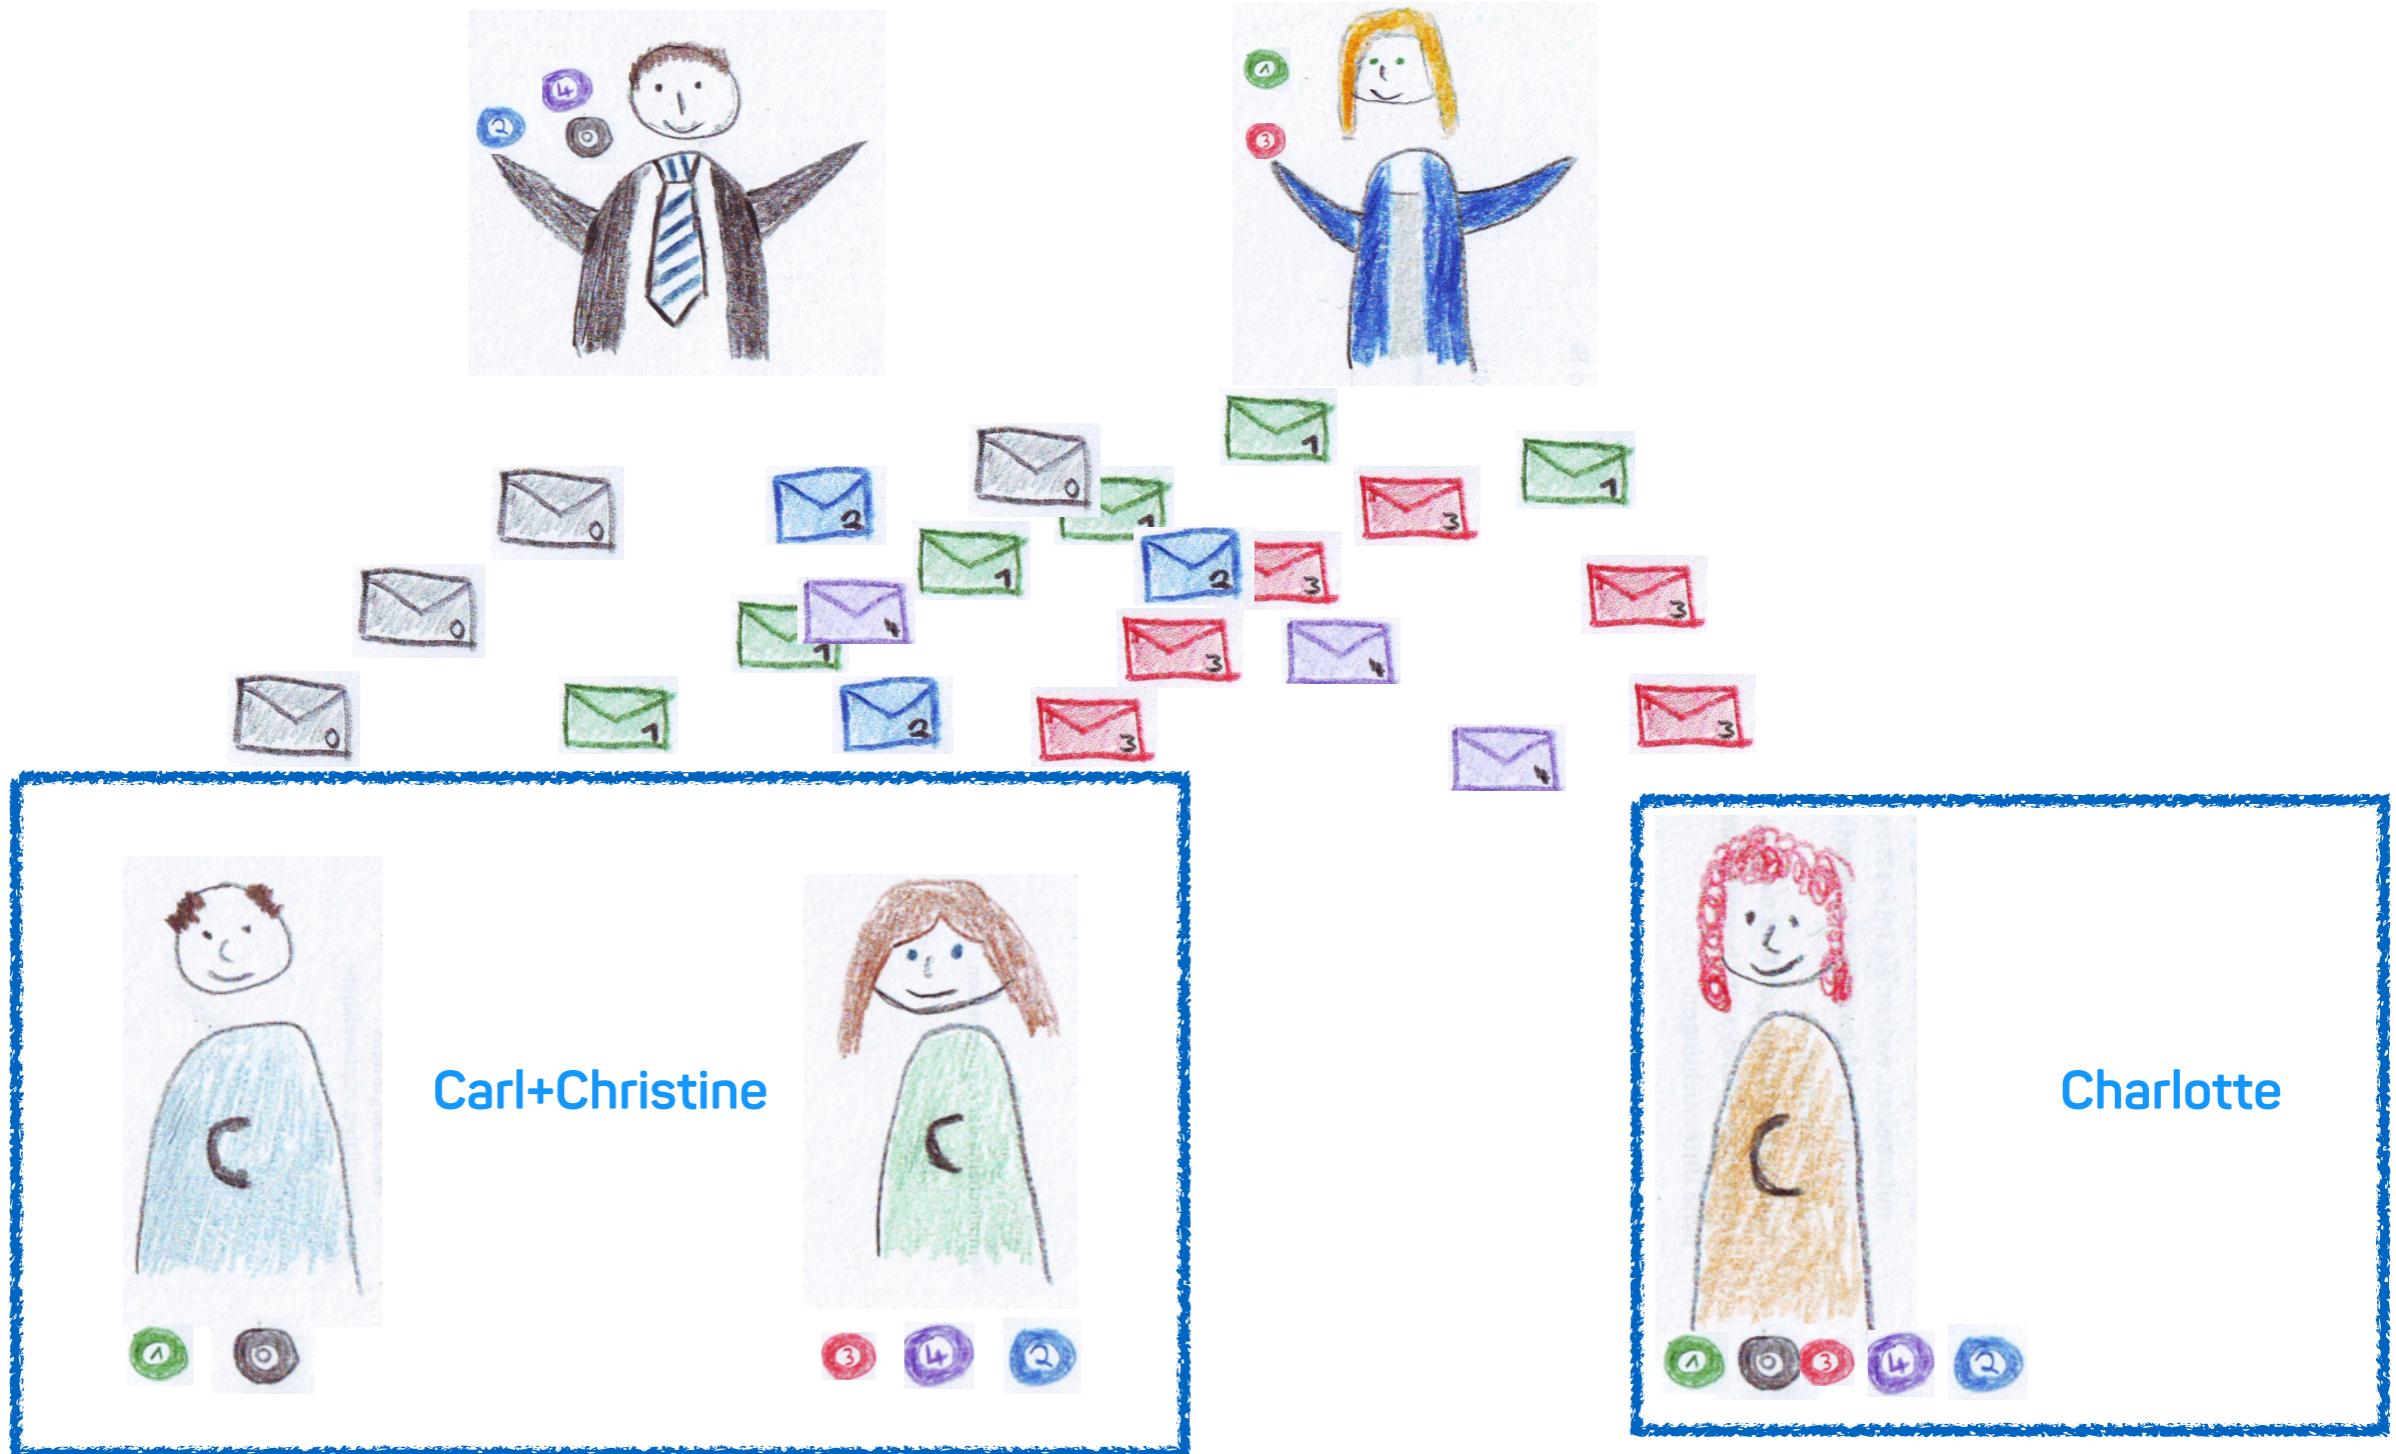
Consumers

- Subscribe to a topic
- Get partitions assigned
- Specify their “Consumer Groups”
 - Partitions are automatically divided among all consumers in a Consumer Group
- Only one concurrent consumer per partition and Consumer Group

Offsets

- Consumers commit their read offsets
 - automatically
 - manually
- Can restart from last committed offset after failure or restart
- Can rewind to historical offsets

Consumers



Failure tolerance

Each topic gets a replication factor

- One broker acts as leader for a partition
- n brokers serve as followers
 - Just serving as backup, never delivering messages for that partition!
- On leader failure, there is a leadership election
 - External dependency on Apache Zookeeper!
 - Leader gets reelected from pool of In-Sync-Replicas
 - Clients get notified via Kafka Protocol

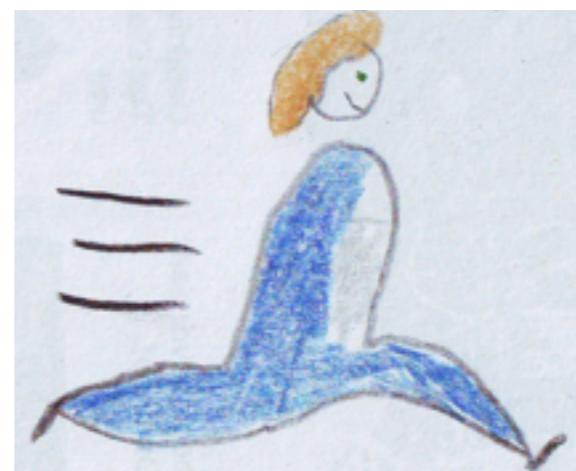
Replication factor 2



Replication as pull by followers



Replicas that are caught up with the leader (per partition)



In-Sync-Replicas

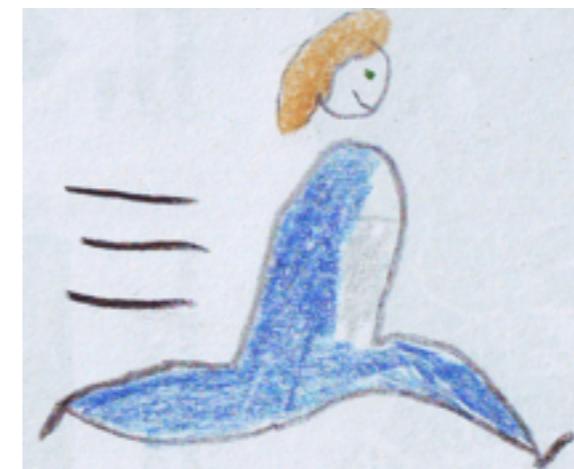


Leader

Replicas that are no longer caught up with the leader



Out-of-Sync-Replica



In-Sync-Replica



Leader

Producers set desired Consistency Level

- Number of acknowledgements from In-Sync-Replicas
- Acknowledgement to Producer after CL serviced
 - 0
 - *no acknowledgement required*
 - 1
 - *as soon as leader writes message*
 - all
 - *when all In-Sync-Replicas answered*
 - *if there is no In-Sync-Replica, this will be satisfied when Leader is done!*